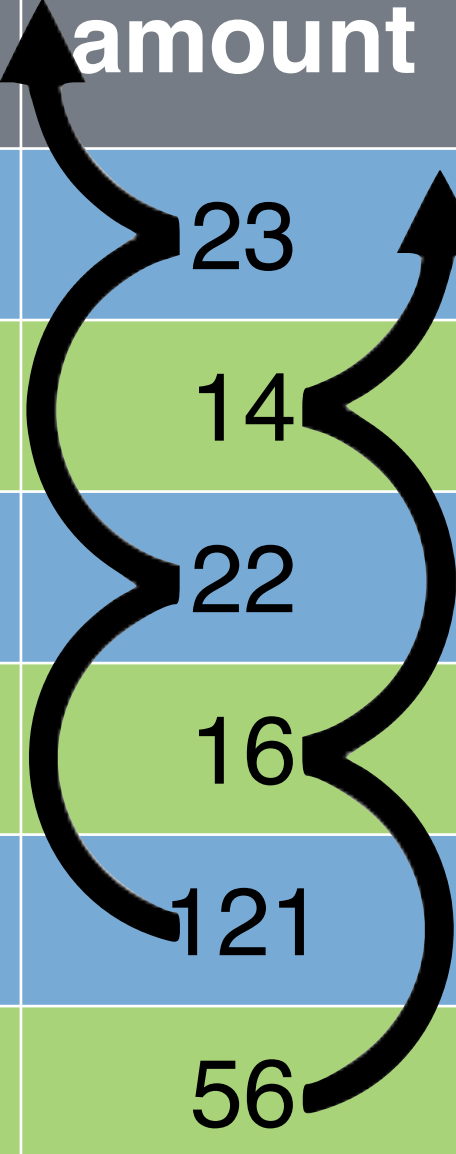


Expert Data Wrangling with R

How to make the most of your data

city	size	amount
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56



Garrett Grolemund

Data Scientist and Master Instructor
November 2014
Email: garrett@rstudio.com

Introduction

```
print.R x DESCRIPTION x html_document.Rmd x handle_click.R x vega.R x
Source on Save Run Source
25 }
26
27 data_props <- combine_data_props(x$marks)
28 data_ids <- names(data_props)
29 data_table <- x$data[data_ids]
30
31 # Collapse each list of scale objects into one scale object.
32 x <- collapse_scales(x)
33 scale_data_table <- scale_domain_data(x)
34
35 # Wrap each of the reactive data objects in another reactive which returns
36 # only the columns that are actually used, and adds any calculated columns
37 # that are used in the props.
38 data_table <- active_props(data_table, data_props)
39
40 # From an environment containing data_table objects, get static data for the
41 # specified ids.
16:19 (Top Level) R Script
```

Environment History Build Git

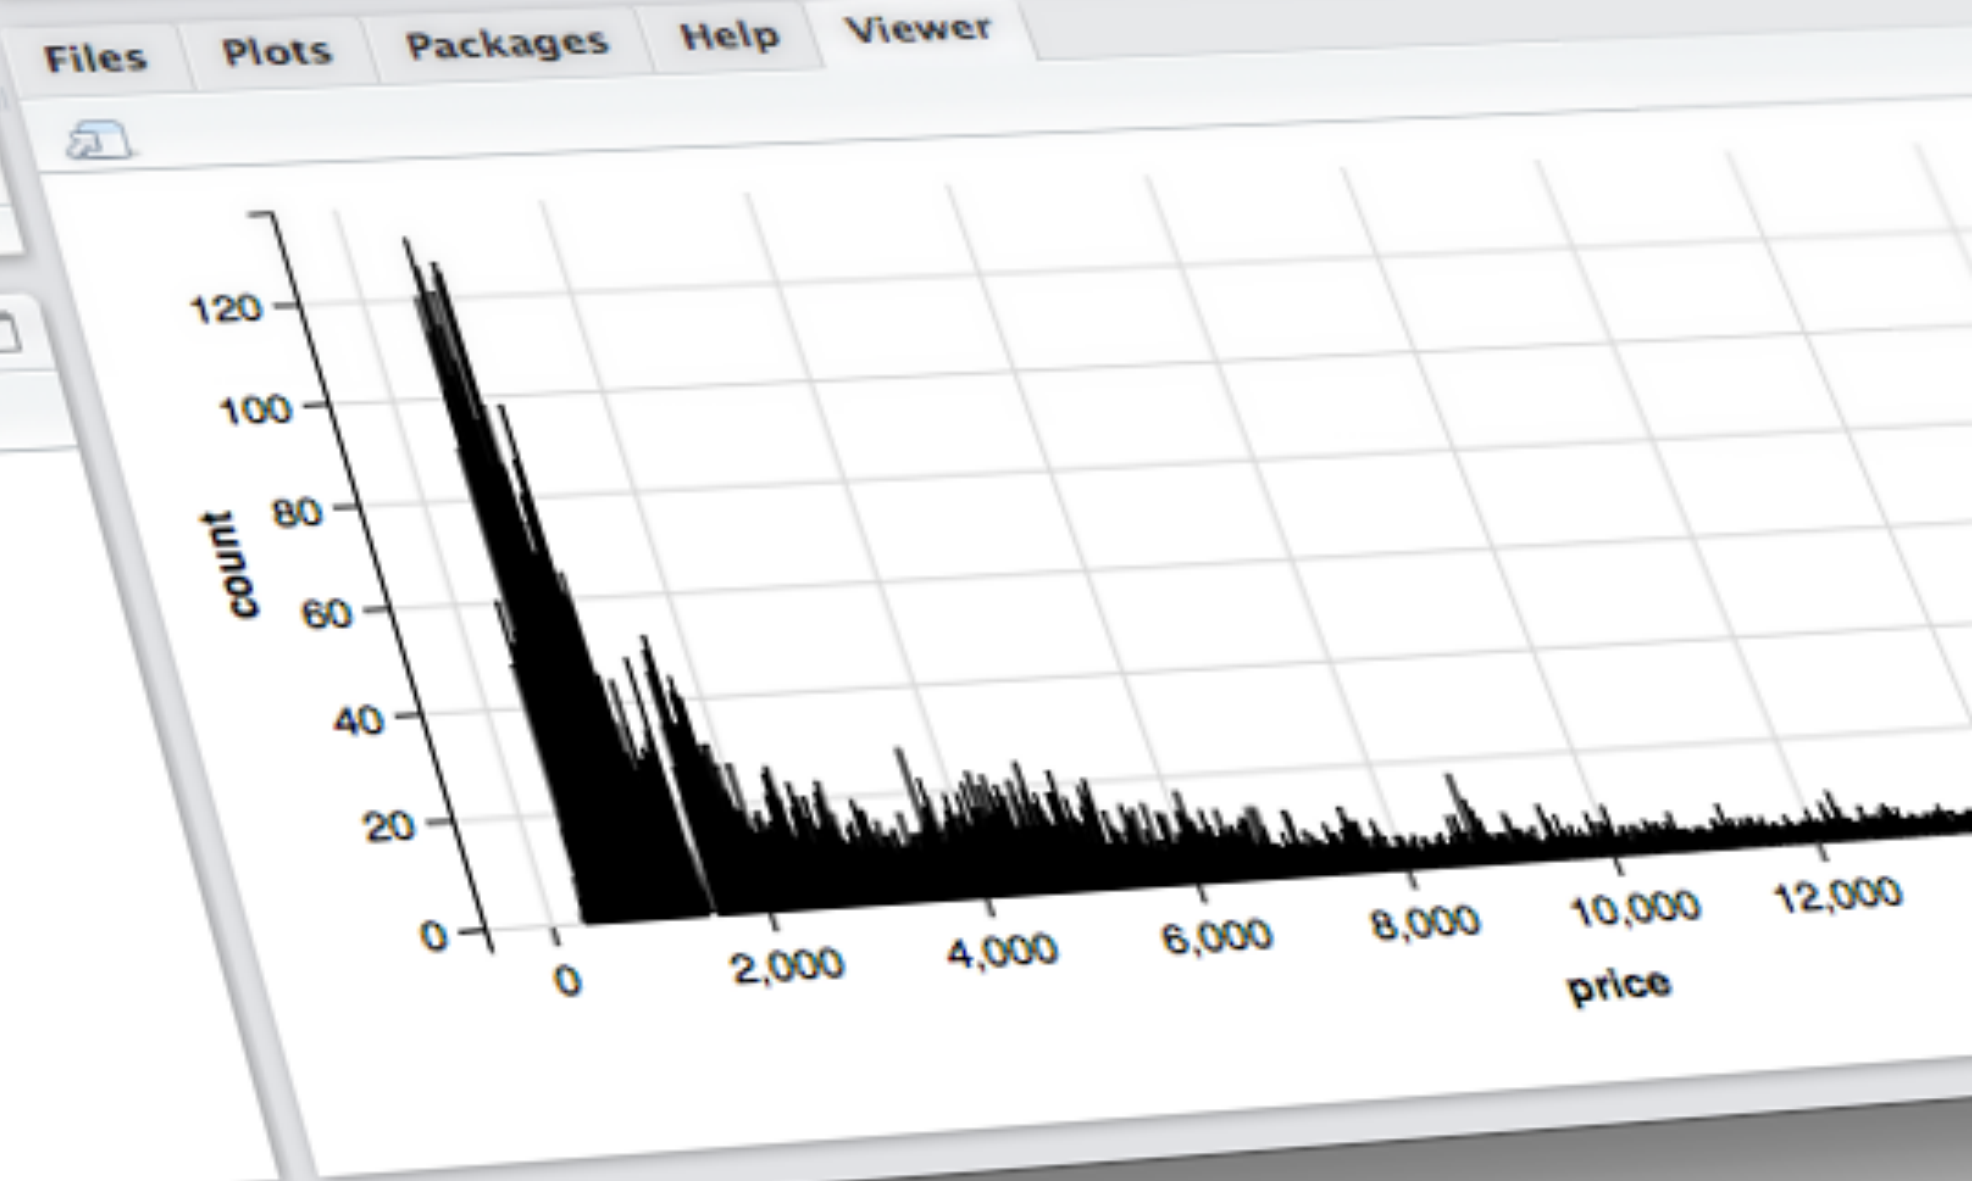
as.vega.ggvis() ~

Values

data_ids	"diamonds0/bin1/stack2"
data_props	List of 1
dynamic	FALSE

Traceback

- as.vega.ggvis(x, FALSE) at vega.R:29
- as.vega(x, FALSE) at vega.R:11
- view_static(x, ...) at print.R:67
- print.ggvis(c("list()", "list(diamonds0 = function () \nstatic_data)", "li



Console ~/r/ggvis/

Next Continue Stop

```
> ggvis(diamonds, x = ~price, y = ~color)
Guessing layer_histograms()
Guessing binwidth = 1
Called from: eval(expr, envir, enclos)
Browse[1]> n
debug at /Users/jmcphers/r/ggvis/R/vega.R#29: data_table <- x$data[data_ids]
Browse[2]>
```


Wrangling

Munging

Cleaning

Carpentry

Manipulation

Transformation

50-80%
of your time?



HELLO

my name is

Garrett



@StatGarrett



O'REILLY®



Hands-On Programming with R

WRITE YOUR OWN FUNCTIONS AND SIMULATIONS

Garrett Grolmund
Foreword by Hadley Wickham

O'REILLY®



Introduction to Data Science with R

Garrett Grolmund

VIDEO

Data Wrangling

Data Wrangling

The practical task of transforming the *layout*, *substance*, and *display* of your data.

Data Manipulation



Changing the variables, values, and units of analysis contained in the data set.

Data Tidying


Changing the layout of tabular data to make it suitable for a particular piece of software (R).

Data Visualization


Transforming the data to a visual format that reveals visual patterns.




Three packages to help you
work with the structure of data.



dplyr



tidyr



ggvis

Data Wrangling with dplyr and tidyr

Cheat Sheet



Syntax - Helpful conventions for wrangling

dplyr::tbl_df(iris)

Converts data to tbl class. tbl's are easier to examine than data frames. R displays only the data that fits onscreen:

```
Source: local data frame [150 x 5]
  Sepal.Length Sepal.Width Petal.Length
1           5.1           3.5           1.4
2           4.9           3.0           1.4
3           4.7           3.2           1.3
4           4.6           3.1           1.5
5           5.0           3.6           1.4
..          ...           ...           ...
Variables not shown: Petal.Width (dbl),
Species (fctr)
```

dplyr::glimpse(iris)

Information dense summary of tbl data.

utils::View(iris)

View data set in spreadsheet-like display (note capital V).

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa
7	4.6	3.4	1.4	0.3	setosa
8	5.0	3.4	1.5	0.2	setosa

dplyr::%>%


Passes object on left hand side as first argument (or argument) of function on righthand side.

`x %>% f(y)` is the same as `f(x, y)`
`y %>% f(x, ., z)` is the same as `f(x, y, z)`

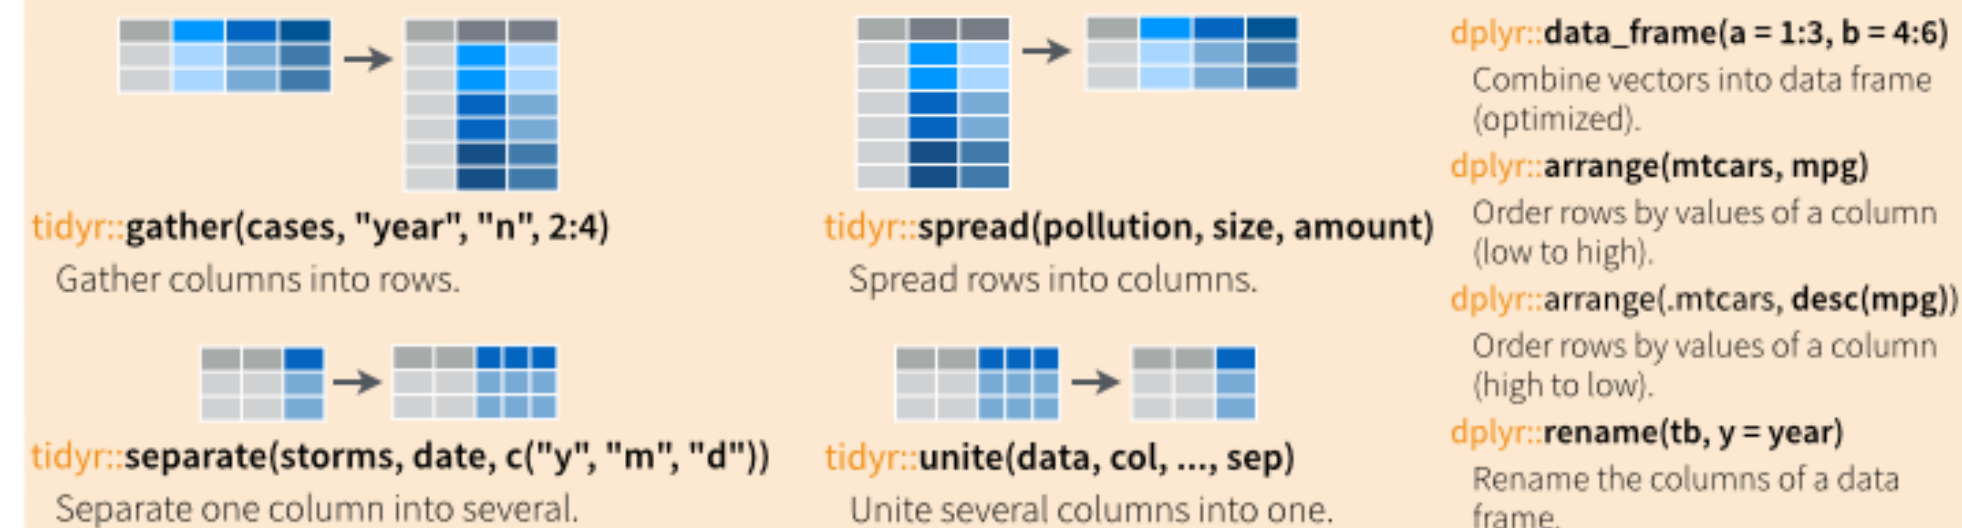
"Piping" with %>% makes code more readable, e.g.

```
iris %>%
  group_by(Species) %>%
  summarise(avg = mean(Sepal.Width)) %>%
  arrange(avg)
```

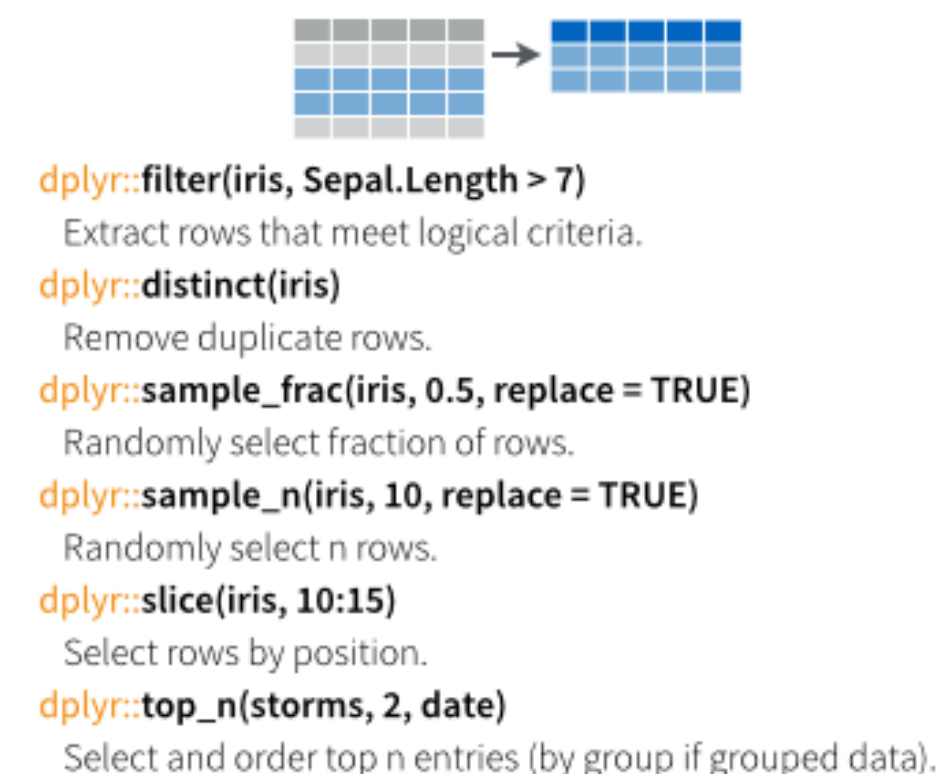
Tidy Data - A foundation for wrangling in R



Reshaping Data - Change the layout of a data set



Subset Observations (Rows)



Subset Variables (Columns)




Helper functions for select - ?select

```
select(iris, contains("."))
  Select columns whose name contains a character string.
select(iris, ends_with("Length"))
  Select columns whose name ends with a character string.
select(iris, everything())
  Select every column.
select(iris, matches(".t."))
  Select columns whose name matches a regular expression.
select(iris, num_range("x", 1:5))
  Select columns named x1, x2, x3, x4, x5.
select(iris, one_of(c("Species", "Genus")))
  Select columns whose names are in a group of names.
select(iris, starts_with("Sepal"))
  Select columns whose name starts with a character string.
select(iris, Sepal.Length:Petal.Width)
  Select all columns between Sepal.Length and Petal.Width (inclusive).
select(iris, -Species)
  Select all columns except Species.
```


Logic in R - ?Comparison, ?base::Logic

<	Less than	!=	Not equal to
>	Greater than	%in%	Group membership
==	Equal to	is.na	Is NA
<=	Less than or equal to	!is.na	Is not NA
>=	Greater than or equal to	&, , !, xor, any, all	Boolean operators

R Packages with example data sets




EDAWR



nycflights13

Case Study

Explore the spread of TB from 1995 to 2013 in 100 countries.



Your Turn

Open R. Then download the packages that we will use today by running the following code.



EDAWR

```
install.packages("devtools")  
devtools::install_github("rstudio/EDAWR")
```




dplyr, tidyr, ggvis, nycflights13


```
install.packages(c("dplyr", "tidyr",  
  "ggvis", "nycflights13"))
```


**Two new
conventions**

Install packages



```
library(dplyr)
```



```
library(nycflights13)
```

nycflights13

Arrival, departure, meterological, etc. data for 300,000+ flights that flew in and out of New York City in 2013.



tbl's

tbl's

Just like data frames, but play better with the console window.

Source: local data frame [336,776 x 16]

	year	month	day	dep_time	dep_delay	arr_time
1	2013	1	1	517	2	830
2	2013	1	1	533	4	850
3	2013	1	1	542	2	923
4	2013	1	1	544	-1	1004
5	2013	1	1	554	-6	812
6	2013	1	1	554	-4	740
7	2013	1	1	555	-5	913
8	2013	1	1	557	-3	709
9	2013	1	1	557	-3	838
10	2013	1	1	558	-2	753
..

Variables not shown: arr_delay (dbl), carrier (chr), tailnum (chr), flight (int), origin (chr), dest (chr), air_time (dbl), distance (dbl), hour (dbl), minute (dbl)

tbl

1977	2013	1	3	847	2
1978	2013	1	3	848	-2
1979	2013	1	3	850	-5
1980	2013	1	3	850	21
1981	2013	1	3	851	24
1982	2013	1	3	851	111
1983	2013	1	3	851	-4
1986	2013	1	3	853	-5
1987	2013	1	3	854	144
1988	2013	1	3	855	-3
1989	2013	1	3	855	-5
1990	2013	1	3	855	-5
1991	2013	1	3	856	-12
1992	2013	1	3	856	36
1993	2013	1	3	857	-3
1994	2013	1	3	857	-3
1995	2013	1	3	857	-3
1996	2013	1	3	858	-2
1997	2013	1	3	858	-2
1998	2013	1	3	859	39
1999	2013	1	3	859	-1
2000	2013	1	3	900	3

[reached getOption("max.print") -- omitted 334776 rows]

data.frame

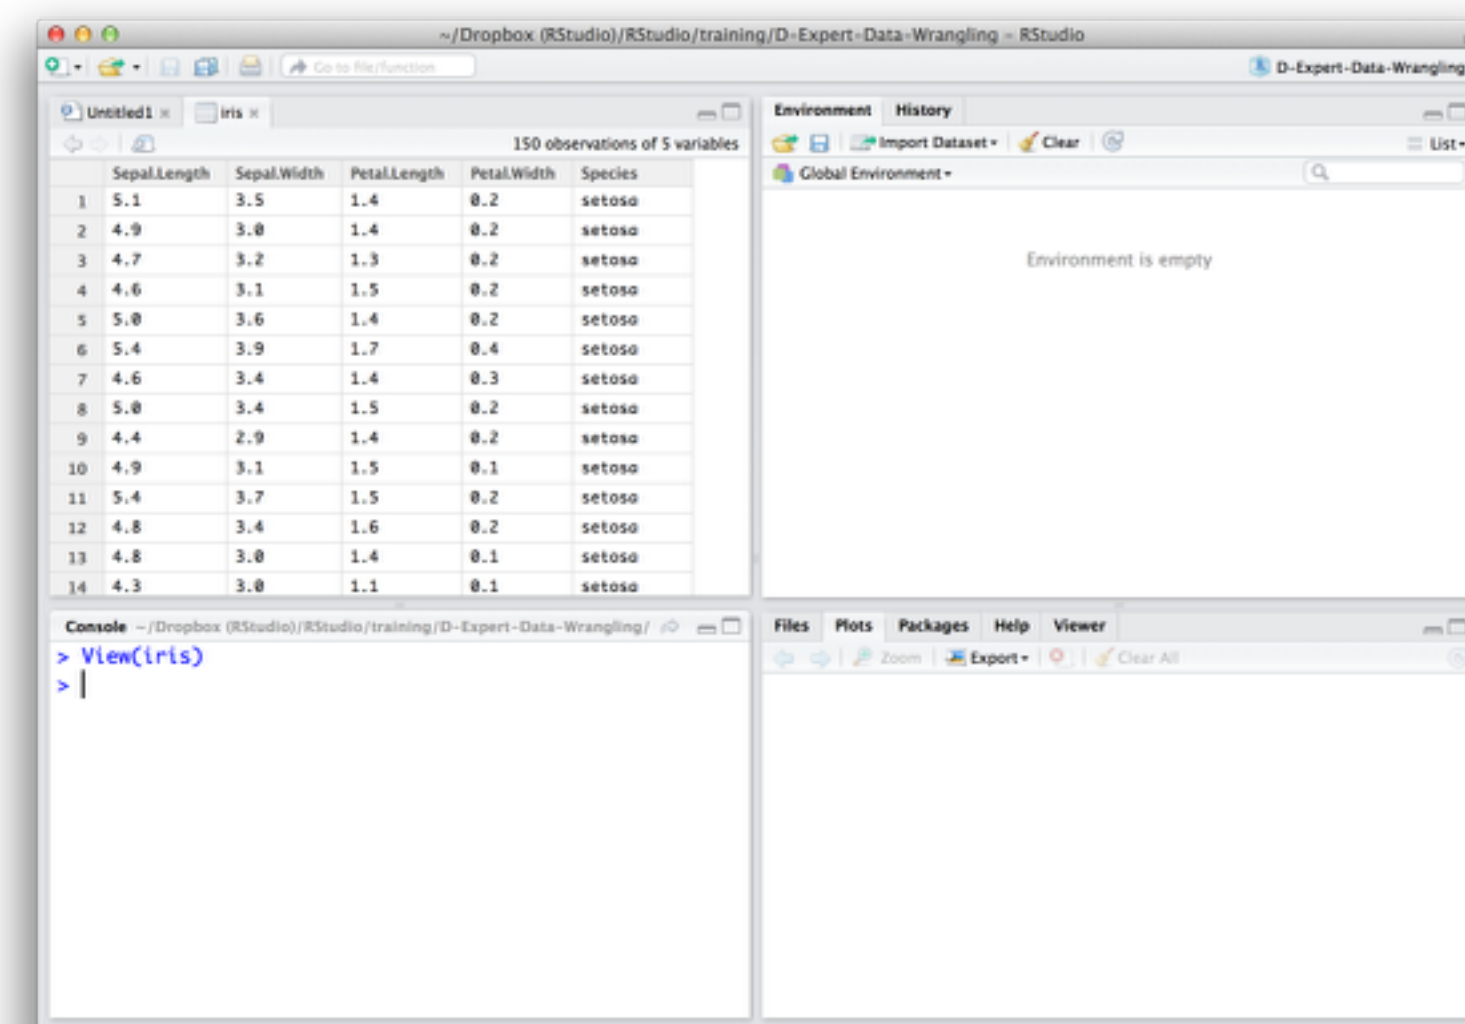
```
flights <- tbl_df(flights)  
# Can undo with  
# flights <- as.data.frame(flights)
```


View()

Examine any data set with the View()
command (Capital V)

```
View(flights)  
View(iris)  
View(mtcars)
```

Data viewer
opens here



glimpse()

Examine values and structure at command line

```
glimpse(flights)  
glimpse(iris)  
glimpse(mtcars)
```

%>%

Ceci n'est pas une pipe.

The pipe operator

%>%

```
dd <- flights$dep_delay  
mean(dd, na.rm = TRUE)  
dd %>% mean(na.rm = TRUE)
```

These do the
same thing

Try it!



dd mean(_____, na.rm = TRUE)


```
little_bunny.foo_foo %>%  
  hop_through(forest) %>%  
  scoop_up(field.mouse) %>%  
  bop_on(head)
```

Data Science*

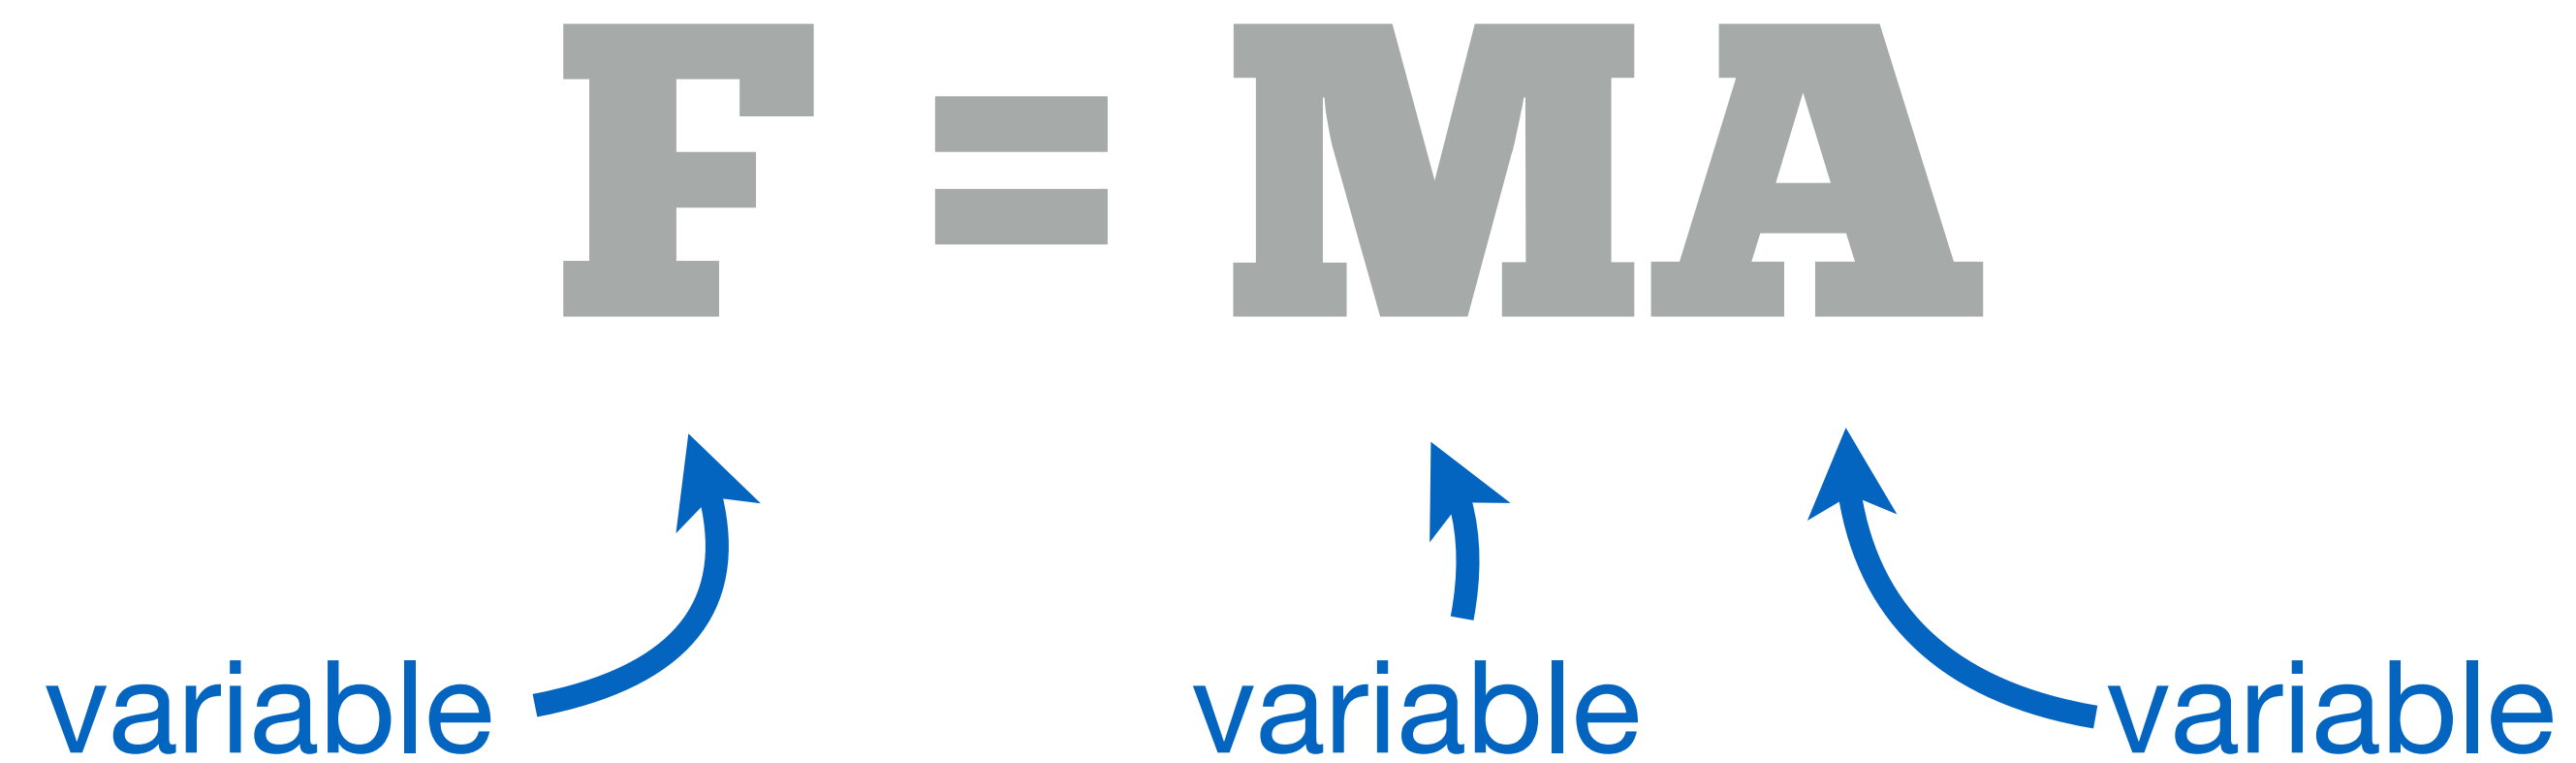
*for Data Wranglers

"Music theory is simply an attempt to explain what musicians have been playing intuitively for thousands of years"

- Mike Iverson

Science

- 1** The best way to learn about the world is to **observe** it.
- 2** Nature behaves according to natural **laws** (e.g. $E = MC^2$, $F = MA$, ...).



Variable - A quantity, quality, or property that you can measure.

$$\mathbf{F} = \mathbf{MA}$$

$$\begin{array}{ccc} f_1 & m_1 & a_1 \\ f_2 & m_2 & a_2 \\ f_3 & m_3 & a_3 \end{array}$$

Variable - A quantity, quality, or property that you can measure.

Value - The state of a variable when you measure it.

(The value can change from measurement to measurement)

$$\mathbf{F} = \mathbf{MA}$$

$$\mathbf{f}_1 = \mathbf{m}_1 \cdot \mathbf{a}_1$$

$$\mathbf{f}_2 = \mathbf{m}_2 \cdot \mathbf{a}_2$$

$$\mathbf{f}_3 = \mathbf{m}_3 \cdot \mathbf{a}_3$$

Variable - A quantity, quality, or property that you can measure.

Value - The state of a variable when you measure it.

Observation - The values of several variables measured under similar conditions.

$$\mathbf{F} = \mathbf{MA}$$

$$f_1 = m_1 \cdot a_1$$


$$f_2 = m_2 \cdot a_2$$

$$f_3 = m_3 \cdot a_3$$

Structure of Natural Laws

Natural laws deal with **variables**,
but they operate on **values**
that appear in the same **observation**.

F	M	A	= MA
f_1	m_1	a_1	$= m_1 \cdot a_1$
f_2	m_2	a_2	$= m_2 \cdot a_2$
f_3	m_3	a_3	$= m_3 \cdot a_3$



- Variables
- Values
- Observations

$$\mathbf{F} = \mathbf{M} \mathbf{A}$$

$$f_1 = m_1 \cdot a_1$$

$$f_2 = m_2 \cdot a_2$$

$$f_3 = m_3 \cdot a_3$$

Data = values associated with
variables and observations

F	M	A
3.01	0.98	3.08
2.35	0.91	2.58
5.57	1.01	5.52

$$\mathbf{F} = \mathbf{MA}$$

$$f_1 = m_1 \cdot a_1$$

$$f_2 = m_2 \cdot a_2$$

$$f_3 = m_3 \cdot a_3$$

Laws appear as
patterns in data.

F	M	A
3.01	0.98	3.08
2.35	0.91	2.58
5.57	1.01	5.52
0.62	1.09	0.56
4.15	0.89	4.69
5.07	1.05	4.84
7.56	0.93	8.12

$$\mathbf{F} = \mathbf{M}\mathbf{A}$$

$$f_1 = m_1 \cdot a_1$$


$$f_2 = m_2 \cdot a_2$$

$$f_3 = m_3 \cdot a_3$$


Laws appear as
patterns in data.

F	M	A
0.62	1.09	0.56
1.30	0.99	1.30
1.63	0.96	1.70
1.72	1.00	1.71
1.82	1.01	1.80
1.95	0.94	2.08
2.11	1.03	2.05

A




F	M	A
0.62	1.09	0.56
1.30	0.99	1.30
1.63	0.96	1.70
1.72	1.00	1.71
1.82	1.01	1.80
1.95	0.94	2.08
2.11	1.03	2.05



Recap: Data science


Structure of Natural Laws

Natural laws deal with **variables**,
but they operate on **values**
that appear in the same **observation**.




Natural laws deal with **variables**, but operate on **values** that appear in the same **observation**.

Data contains **values** grouped into **variables** and **observations**.



Laws appear as patterns in data.

EDAWR



An R package with toy data sets that we will use today.

```
# install.packages("devtools")  
# devtools::install_github("rstudio/EDAWR")  
library(EDAWR)  
?tb  
?population
```

Warm up - Identify the variables

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21

cases

Country	2011	2012	2013
FR	7000	6900	7000
DE	5800	6000	6200
US	15000	14000	13000

pollution

city	particle size	amount ($\mu\text{g}/\text{m}^3$)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ama	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arnold	45	1010	1996-06-21

- Storm name
- Wind Speed (mph)
- Air Pressure
- Date

cases

Country	2011	2012	2013
FR	7000	6900	7000
DE	800	6000	6200
US	15000	12000	13000

- Country
- Year
- Count

pollution

city	particle size	amount (µg/m³)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56

- City
- Amount of large particles
- Amount of small particles

Your Turn - Identify the observations

storms

storm	wind	pressure	date
Alberto	110	1007	2000-08-12
Alex	45	1009	1998-07-30
Allison	65	1005	1995-06-04
Ana	40	1013	1997-07-01
Arlene	50	1010	1999-06-13
Arthur	45	1010	1996-06-21

cases

Country	2011	2012	2013
FR	7000	6900	7000
DE	5800	6000	6200
US	15000	14000	13000

pollution

city	particle size	amount ($\mu\text{g}/\text{m}^3$)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56

storms

storm	wind	pressure	date
Albino	110	1007	2000-09-12
Andrea	45	1005	1998-09-30
Alison	65	1005	1995-09-04
Andrea	40	1010	1997-09-01
Andrea	50	1010	1999-09-13
Andrea	45	1010	1998-09-21

(per storm)

cases

Country	2011	2012	2013
USA	7000	3000	7000
UK	5000	3000	6200
France	15000	14000	3000

(per country per year)

pollution

city	particle size	amount (µg/m³)
New York	large	23
New York	small	14
London	large	22
London	small	16
Beijing	large	121
Beijing	small	56

(per city)